

ReplicaSet:

A ReplicaSet is a Kubernetes Object responsible for ensuring that a specified number of pod replicas are running and operational at any given time. It helps to achieve high availability and scalability for your applications by managing the lifecycle of pods.

The role of the ReplicaSet is to Monitor pods (by matching given labels) if any of them failed then deploy a new one.

Key Concepts:

Desired State: ReplicaSets maintain a desired state, which defines the number of pod replicas that should be running. This desired state is specified in the ReplicaSet configuration.

Pod Template: ReplicaSets use a pod template to create and manage pods. The pod template includes specifications for the pods that the ReplicaSet should create and manage.

Selectors: ReplicaSets use label selectors to identify the pods that they should manage. Pods managed by a ReplicaSet must have labels that match the ReplicaSet's selector.

Pod Creation and Deletion: ReplicaSets continuously monitor the cluster and create or delete pods as necessary to ensure that the actual state matches the desired state.

ReplicaSet Features:

Defining a ReplicaSet: To create a ReplicaSet, you define a YAML or JSON configuration file that specifies the desired number of replicas, the pod template, and any other configuration options.

Creating Pods: When you apply the ReplicaSet configuration, Kubernetes creates the specified number of pod replicas using the pod template.

Monitoring Pods: The ReplicaSet continuously monitors the cluster to ensure that the actual number of pod replicas matches the desired number specified in the configuration.

Scaling: If the actual number of pod replicas deviates from the desired state (e.g., due to pod failures or manual scaling actions), the ReplicaSet takes corrective actions by creating or deleting pods to bring the cluster back to the desired state.

Updating Pods: When you update the pod template in the ReplicaSet configuration (e.g., to deploy a new version of your application), the ReplicaSet creates new pods with the updated template and gradually replaces the existing pods, ensuring zero-downtime deployments.

REPLICASET DEFINITION FILE

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: example-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-container
          image: nginx:latest
```

Commands:

- `kubectl get replicaset` (We can use replicaset or rs)
- `kubectl get rs -A`
- `kubectl describe rs [replicaset name]`
- `kubectl create -f rs-definition.yaml`
- `kubectl apply -f rs-definition.yaml`
- `kubectl get replicaset [replicaset name] -o yaml > file.yaml`
- `kubectl scale replicaset [replicaset name] --replicas [number]`

IMPORTANT

Replicaset vs ReplicationController

Both share a similar core functionality: ensuring a desired number of pod replicas are running at any given time. However, there are some key differences between them:

Selector: This specifies which pods the controller manages. Replica sets use a more powerful set-based selector, allowing for richer matching options like "in", "notIn", "exists", or "doesNotExist" operators for pod labels. Replication controllers, on the other hand, are limited to equality-based selectors.

Updating Pods: Replica sets are primarily designed to be used as the backend for Deployments, which handle rolling updates of pods. Replication controllers have a built-in rollingUpdate command, but it's generally considered less flexible.

Status: Replication controllers are considered deprecated and should be replaced with replica sets for new deployments. Deployments are the recommended approach for managing pod replicas due to their additional features for controlled rollouts and rollbacks.

Feature	ReplicaSet	ReplicationController
Selector	Set-based (More Powerful)	Equality Based
Updating Pods	Designed for use with Deployments	Build-in rollingUpdate Command
Status	Recommended	Deprecated

ReplicationController Definition File

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: my-rc
  labels:
    type: rc
spec:
  replicas: 3
  selector:
    tier: backend
  template:
    metadata:
      name: myapp
      labels:
        tier: backend
    spec:
      containers:
        - name: myapp
          image: nginx:latest
          ports:
            - containerPort: 80
```

ReplicationController Commands

- `kubectl create -f rc-definition.yaml`
- `kubectl get rc`
- `kubectl get rc -o wide`
- `kubectl get rc rc-name -o yaml`
- `kubectl describe rc rc-name`
- `kubectl delete rc rc-name`
- `kubectl rolling-update rc-name --image updatedImage`

While these features provide basic pod management, it's important to consider the limitations of replication controllers:

Limited Rolling Updates: Compared to deployments, replication controller rolling updates are less sophisticated. They lack features like health checks and easy rollbacks in case of issues with the update.

Deprecated Status: Replication controllers are considered deprecated in favor of replica sets with deployments. Deployments offer a more robust and feature-rich approach for managing pod lifecycles and updates.